

1a) Implementing TISP variable selection for classification.

As always start with importing what we need and finding our data locations.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, roc_curve, auc

import os

os.chdir("C:/Users/rique/Downloads/gisette")
```

For part a) We're using the Gisette data

load the data.

```
g_train_data = np.loadtxt("gisette_train.data")
g_train_labels = np.loadtxt("gisette_train.labels")
g_valid_data = np.loadtxt("gisette_valid.data")
g_valid_labels = np.loadtxt("gisette_valid.labels")
```

Normalize the features. Copying and pasting from my HW3, cause it works. Don't fix what ain't broken.

```
std=np.std(g_train_data, axis=0)

#set a mask so we dont get a divide by standard deviation of zero
mask = (std != 0)

#apply the mask, get the mean and standard dev of the normalized data
g_train_data = g_train_data[:, mask]
mean = np.mean(g_train_data, axis=0)
true_std = np.std(g_train_data, axis=0)

#standardize the features in both training and test datasets
g_train_data = (g_train_data-mean)/true_std
g_valid_data = g_valid_data[:, mask]
g_valid_data = (g_valid_data-mean)/true_std

#add a bias term
g_train_data = np.insert(g_train_data, 0, 1, axis=1)
g_valid_data = np.insert(g_valid_data, 0, 1, axis=1)
```

```
g_train_labels[g_train_labels == 0] = -1
g_valid_labels[g_valid_labels == 0] = -1
```

Verifying data has mean 0 and variance of 1

```
print("Train mean: ", np.mean(g_train_data))
print("Train variance: ", np.var(g_train_data))
print("Test mean: ", np.mean(g_valid_data))
print("Test variance: ", np.var(g_valid_data))
```

```
Train mean:  0.0002017756255044388
Train variance:  0.9999999592865941
Test mean:  0.0062998654925201245
Test variance:  1.0634350583552006
```

Initialize our needed parameters

threshold values are subject to change.... ALOT.

```
iterations = 100

#lambda; these values are subject to change to find out features

#For lambda of: 0.087959 feature is: 98
#0.08795105 101

#For lambda of: 0.038549 feature is: 499
#For lambda of: 0.038535 feature is: 502
#For lambda of: 0.038541 feature is: 502

thresholds = [0.19, 0.133, 0.08795105, 0.05291, 0.038545]

w = np.zeros(g_train_data.shape[1])

train_misclass_errors = []
valid_misclass_errors = []
train_misclass_errors_30 = []
features = []
fpr_train_list = []
tpr_train_list = []
roc_auc_train_list = []
fpr_valid_list = []
tpr_valid_list = []
roc_auc_valid_list = []
```

TISP implementation

```
for lambda_ in thresholds:
    for i in range(iterations):
        # Dot product of train data and weight
        dot = np.sum(g_train_data * w, axis=1)

        # Gradient
        gradient = np.sum((g_train_labels / (1 + np.exp(g_train_labels
* dot))) * (g_train_data).T, axis=1)

        # Update the weight with our gradient
        w += gradient * (1 / g_train_data.shape[0])
        w[np.absolute(w) <= lambda_] = 0

        # print(i, "weight is: ", np.sum(w != 0))

        # Recalculate dot product of train data and updated weight
        dot = np.sum(g_train_data * w, axis=1)

        # Prediction here is based on if the dot product of train/test
sets is greater than zero
        y_pred_train = ((dot >= 0) == g_train_labels)
        misclass_error_train = 1 - accuracy_score(g_train_labels,
y_pred_train)

        #cause we wanna plot for 30 features
        if(lambda_ == 0.133):
            train_misclass_errors_30.append(misclass_error_train)

        feature = np.sum(w != 0)
        features.append(feature)
        # print("For lambda of:", lambda_, "feature is:", feature)

        train_misclass_errors.append(misclass_error_train)

        dot_valid = np.sum(g_valid_data * w, axis=1)
        y_pred_valid = ((dot_valid >= 0) == g_valid_labels)
        misclass_error_valid = 1 - accuracy_score(g_valid_labels,
y_pred_valid)
        valid_misclass_errors.append(misclass_error_valid)

        # Calculate ROC curve values for the training set
        fpr_train, tpr_train, _ = roc_curve(g_train_labels, 1 / (1 +
np.exp(-dot)))
        roc_auc_train = auc(fpr_train, tpr_train)
        fpr_train_list.append(fpr_train)
        tpr_train_list.append(tpr_train)
        roc_auc_train_list.append(roc_auc_train)

        # Calculate ROC curve values for the validation set
```

```

fpr_valid, tpr_valid, _ = roc_curve(g_valid_labels, 1 / (1 +
np.exp(-dot_valid)))
roc_auc_valid = auc(fpr_valid, tpr_valid)
fpr_valid_list.append(fpr_valid)
tpr_valid_list.append(tpr_valid)
roc_auc_valid_list.append(roc_auc_valid)

w=np.zeros_like(w)

#print("Features selected:", features)
#print("Train misclassification errors:", train_misclass_errors)
#print("Validation misclassification errors:", valid_misclass_errors)

```

Plot the stuff

```

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(orange(iterations), train_misclass_errors_30, label="Train")
plt.xlabel('Iterations')
plt.ylabel('Miss Class Error')
plt.title('30 Feature: Iterations vs Miss Classification Error')
plt.grid()
plt.legend()

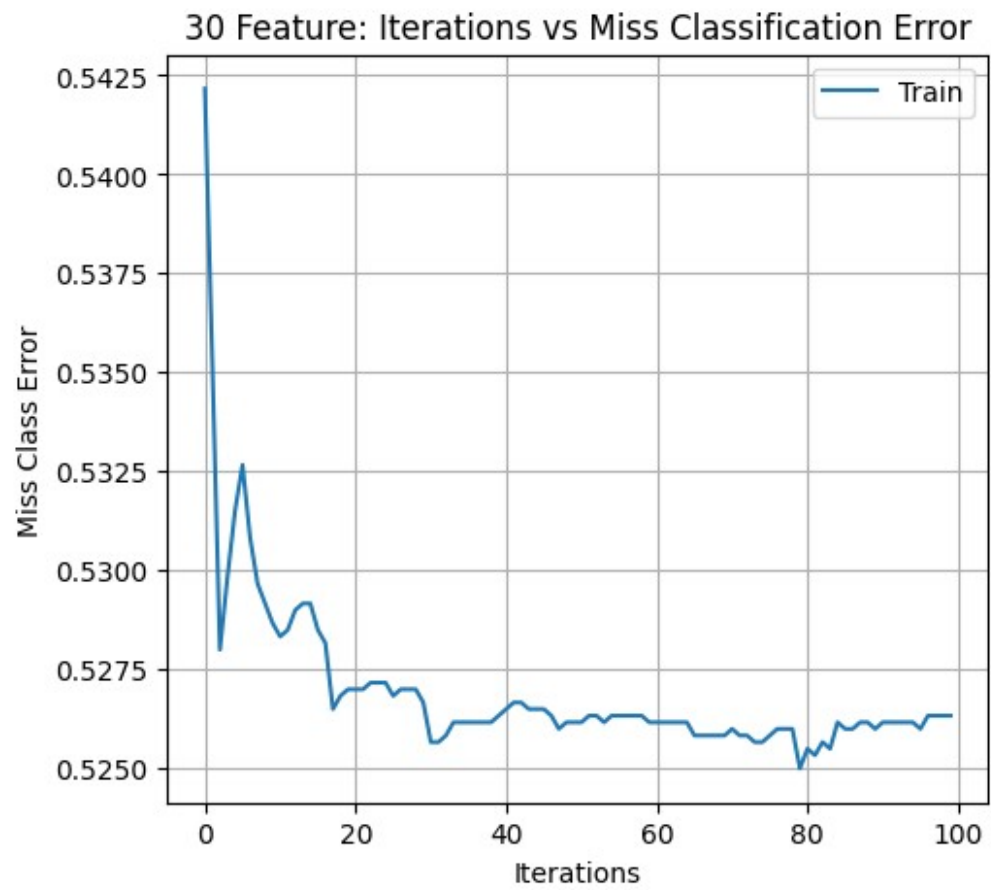
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(features, train_misclass_errors, marker="o", label="Train")
plt.plot(features, valid_misclass_errors, marker="o", label="Test")
plt.xlabel('Features')
plt.ylabel('Miss Class Error')
plt.title('Selected Features vs Miss Classification Error')
plt.grid()
plt.legend()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 2)
plt.plot(fpr_train_list[-1], tpr_train_list[-1], color='blue', lw=2,
label=f'Training Set (AUC = {roc_auc_train_list[-1]:.2f})')
plt.plot(fpr_valid_list[-1], tpr_valid_list[-1], color='darkorange',
lw=2, label=f'Validation Set (AUC = {roc_auc_valid_list[-1]:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='No
Discrimination')


plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()

plt.show()

```



Selected Features vs Miss Classification Error






Table with the Misclass errors, lambda, features on test and training sets

```
results = pd.DataFrame({
    'Lambda': thresholds,
    'Features': features,
    'Train MisClass Error': train_misclass_errors,
    'Test MisClass Error': valid_misclass_errors
})
```

```
print(results)
```

	Lambda	Features	Train MisClass Error	Test MisClass Error
0	0.190000	10	0.545167	0.545
1	0.133000	30	0.526333	0.529
2	0.087951	101	0.514333	0.517
3	0.052910	300	0.505167	0.506
4	0.038545	501	0.504000	0.506

1b) Repeating a) but on the dexter set.

Copy & Pasting above code, with only differences being the dataset, and another guessing game with the thresholds

```
os.chdir("C:/Users/rique/Downloads/dexter")

X = np.genfromtxt('dexter_train.csv', delimiter=',')
y = np.loadtxt('dexter_train.labels')

Xtest = np.genfromtxt('dexter_valid.csv', delimiter=',')
yTest = np.loadtxt('dexter_valid.labels')

std=np.std(X, axis=0)

#set a mask so we dont get a divide by standard deviation of zero
mask = (std != 0)

#apply the mask, get the mean and standard dev of the normalized data
X = X[:, mask]
mean = np.mean(X, axis=0)
true_std = np.std(X, axis=0)

#standardize the features in both training and test datasets
X = (X-mean)/true_std
Xtest = Xtest[:, mask]
Xtest = (Xtest-mean)/true_std

#add a bias term
X = np.insert(X, 0, 1, axis=1)
Xtest = np.insert(Xtest, 0, 1, axis=1)

y[y == 0] = -1
yTest[yTest == 0] = -1

iterations = 100

#lambda; these values are subject to change to find out features

thresholds = [0.14, 0.098, 0.071, 0.0525, 0.0468]

w = np.zeros(X.shape[1])

train_misclass_errors = []
valid_misclass_errors = []
train_misclass_errors_30 = []
features = []
fpr_train_list = []
tpr_train_list = []
roc_auc_train_list = []
```



```

fpr_valid_list = []
tpr_valid_list = []
roc_auc_valid_list = []

for lambda_ in thresholds:
    for i in range(iterations):
        # Dot product of train data and weight
        dot = np.sum(X * w, axis=1)

        # Gradient
        gradient = np.sum((y / (1 + np.exp(y * dot))) * (X).T, axis=1)

        # Update the weight with our gradient
        w += gradient * (1 / X.shape[0])
        w[np.absolute(w) <= lambda_] = 0

        #print(i, "weight is: ", np.sum(w != 0))

        # Recalculate dot product of train data and updated weight
        dot = np.sum(X * w, axis=1)

        # Prediction here is based on if the dot product of train/test
sets is greater than zero
        y_pred_train = ((dot >= 0) == y)
        misclass_error_train = 1 - accuracy_score(y, y_pred_train)

        #cause we wanna plot for 30 features
        if(lambda_ == 0.098):
            train_misclass_errors_30.append(misclass_error_train)

        feature = np.sum(w != 0)
        features.append(feature)
    # print("For lambda of:", lambda_, "feature is:", feature)

    train_misclass_errors.append(misclass_error_train)

    dot_valid = np.sum(Xtest * w, axis=1)
    y_pred_valid = ((dot_valid >= 0) == yTest)
    misclass_error_valid = 1 - accuracy_score(yTest, y_pred_valid)
    valid_misclass_errors.append(misclass_error_valid)

    # Calculate ROC curve values for the training set
    fpr_train, tpr_train, _ = roc_curve(y, 1 / (1 +
np.exp(-dot)))
    roc_auc_train = auc(fpr_train, tpr_train)
    fpr_train_list.append(fpr_train)
    tpr_train_list.append(tpr_train)
    roc_auc_train_list.append(roc_auc_train)

    # Calculate ROC curve values for the validation set

```

```

fpr_valid, tpr_valid, _ = roc_curve(yTest, 1 / (1 +
np.exp(-dot_valid)))
roc_auc_valid = auc(fpr_valid, tpr_valid)
fpr_valid_list.append(fpr_valid)
tpr_valid_list.append(tpr_valid)
roc_auc_valid_list.append(roc_auc_valid)

w=np.zeros_like(w)

#print("Features selected:", features)
#print("Train misclassification errors:", train_misclass_errors)
#print("Validation misclassification errors:", valid_misclass_errors)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(orange(iterations), train_misclass_errors_30, label="Train")
plt.xlabel('Iterations')
plt.ylabel('Miss Class Error')
plt.title('30 Feature: Iterations vs Miss Classification Error')
plt.grid()
plt.legend()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(features, train_misclass_errors, marker="o", label="Train")
plt.plot(features, valid_misclass_errors, marker="o", label="Test")
plt.xlabel('Features')
plt.ylabel('Miss Class Error')
plt.title('Selected Features vs Miss Classification Error')
plt.grid()
plt.legend()


plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 2)
plt.plot(fpr_train_list[-1], tpr_train_list[-1], color='blue', lw=2,
label=f'Training Set (AUC = {roc_auc_train_list[-1]:.2f})')
plt.plot(fpr_valid_list[-1], tpr_valid_list[-1], color='darkorange',
lw=2, label=f'Validation Set (AUC = {roc_auc_valid_list[-1]:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='No
Discrimination')


plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()


plt.show()

```

30 Feature: Iterations vs Miss Classification Error







```
results = pd.DataFrame({  
    'Lambda': thresholds,  
    'Features': features,  
    'Train MisClass Error': train_misclass_errors,  
    'Test MisClass Error': valid_misclass_errors  
})
```

```
print(results)
```

	Lambda	Features	Train MisClass Error	Test MisClass Error
0	0.1400	11	0.596667	0.616667
1	0.0980	30	0.530000	0.556667
2	0.0710	101	0.510000	0.533333
3	0.0525	294	0.500000	0.513333
4	0.0468	495	0.500000	0.523333

1c) Repeating a) and b) but on the madelon set.

Copy & Pasting above code, with only differences being the dataset, and yet another tedious guessing game with the thresholds

```
os.chdir("C:/Users/rique/Downloads/MADELON")

X = np.loadtxt("madelon_train.data")
y = np.loadtxt("madelon_train.labels")

Xtest = np.loadtxt("madelon_valid.data")
yTest = np.loadtxt("madelon_valid.labels")

std=np.std(X, axis=0)

#set a mask so we dont get a divide by standard deviation of zero
mask = (std != 0)

#apply the mask, get the mean and standard dev of the normalized data
X = X[:, mask]
mean = np.mean(X, axis=0)
true_std = np.std(X, axis=0)

#standardize the features in both training and test datasets
X = (X-mean)/true_std
Xtest = Xtest[:, mask]
Xtest = (Xtest-mean)/true_std

#add a bias term
X = np.insert(X, 0, 1, axis=1)
Xtest = np.insert(Xtest, 0, 1, axis=1)

y[y == 0] = -1
yTest[yTest == 0] = -1

iterations = 100

#lambda; these values are subject to change to find out features
#For lambda of: 0.0009 feature is: 473

thresholds = [0.02999, 0.0245, 0.017, 0.0075, 0.000199]

w = np.zeros(X.shape[1])

train_misclass_errors = []
valid_misclass_errors = []
train_misclass_errors_30 = []
features = []
```

```

fpr_train_list = []
tpr_train_list = []
roc_auc_train_list = []
fpr_valid_list = []
tpr_valid_list = []
roc_auc_valid_list = []

for lambda_ in thresholds:
    for i in range(iterations):
        # Dot product of train data and weight
        dot = np.sum(X * w, axis=1)

        # Gradient
        gradient = np.sum((y / (1 + np.exp(y * dot))) * (X).T, axis=1)

        # Update the weight with our gradient
        w += gradient * (1 / X.shape[0])
        w[np.absolute(w) <= lambda_] = 0

        # print(i, "weight is: ", np.sum(w != 0))

        # Recalculate dot product of train data and updated weight
        dot = np.sum(X * w, axis=1)

        # Prediction here is based on if the dot product of train/test
sets is greater than zero
        y_pred_train = ((dot >= 0) == y)
        misclass_error_train = 1 - accuracy_score(y, y_pred_train)

        #cause we wanna plot for 30 features
        if(lambda_ == 0.0245):
            train_misclass_errors_30.append(misclass_error_train)

        feature = np.sum(w != 0)
        features.append(feature)
        #print("For lambda of:", lambda_, "feature is:", feature)

        train_misclass_errors.append(misclass_error_train)

        dot_valid = np.sum(Xtest * w, axis=1)
        y_pred_valid = ((dot_valid >= 0) == yTest)
        misclass_error_valid = 1 - accuracy_score(yTest, y_pred_valid)
        valid_misclass_errors.append(misclass_error_valid)

        # Calculate ROC curve values for the training set
        fpr_train, tpr_train, _ = roc_curve(y, 1 / (1 +
np.exp(-dot)))
        roc_auc_train = auc(fpr_train, tpr_train)
        fpr_train_list.append(fpr_train)
        tpr_train_list.append(tpr_train)

```

```

roc_auc_train_list.append(roc_auc_train)

# Calculate ROC curve values for the validation set
fpr_valid, tpr_valid, _ = roc_curve(yTest, 1 / (1 +
np.exp(-dot_valid)))
roc_auc_valid = auc(fpr_valid, tpr_valid)
fpr_valid_list.append(fpr_valid)
tpr_valid_list.append(tpr_valid)
roc_auc_valid_list.append(roc_auc_valid)

w=np.zeros_like(w)

#print("Features selected:", features)
#print("Train misclassification errors:", train_misclass_errors)
#print("Validation misclassification errors:", valid_misclass_errors)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(orange(iterations), train_misclass_errors_30, label="Train")
plt.xlabel('Iterations')
plt.ylabel('Miss Class Error')
plt.title('30 Feature: Iterations vs Miss Classification Error')
plt.grid()
plt.legend()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(features, train_misclass_errors, marker="o", label="Train")
plt.plot(features, valid_misclass_errors, marker="o", label="Test")
plt.xlabel('Features')
plt.ylabel('Miss Class Error')
plt.title('Selected Features vs Miss Classification Error')
plt.grid()
plt.legend()

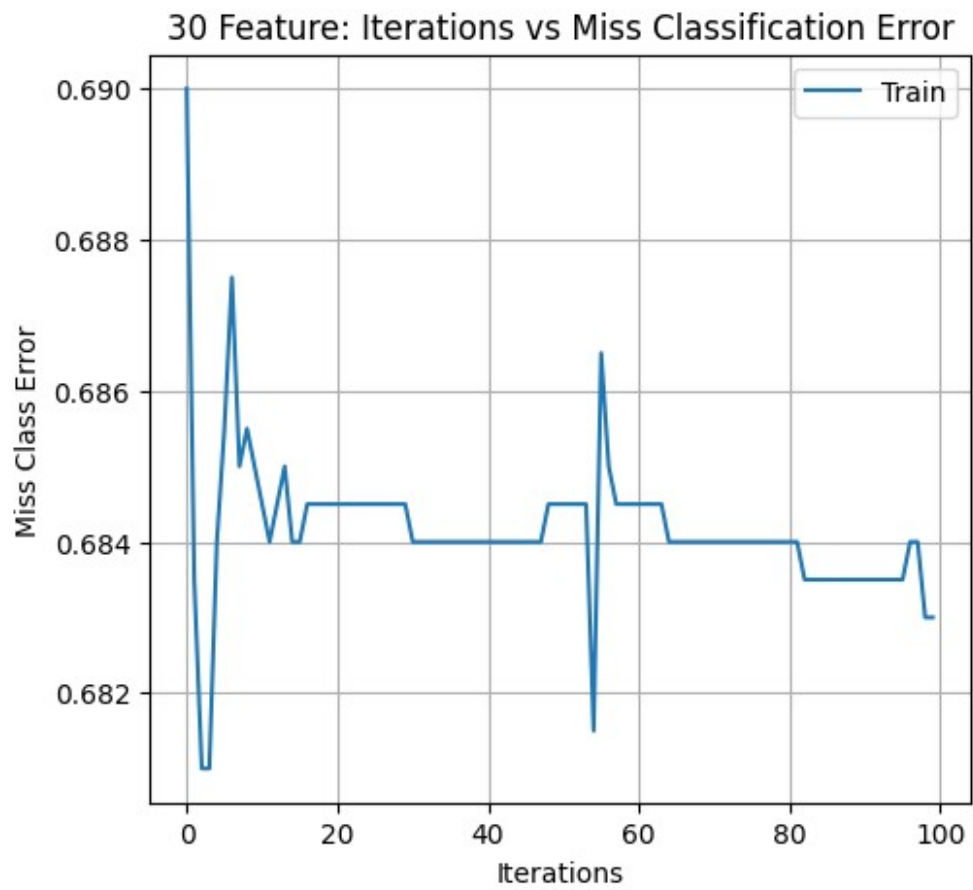
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 2)
plt.plot(fpr_train_list[-1], tpr_train_list[-1], color='blue', lw=2,
label=f'Training Set (AUC = {roc_auc_train_list[-1]:.2f})')
plt.plot(fpr_valid_list[-1], tpr_valid_list[-1], color='darkorange',
lw=2, label=f'Validation Set (AUC = {roc_auc_valid_list[-1]:.2f})')
# Add a dashed diagonal line (no-discrimination ROC curve)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='No
Discrimination')

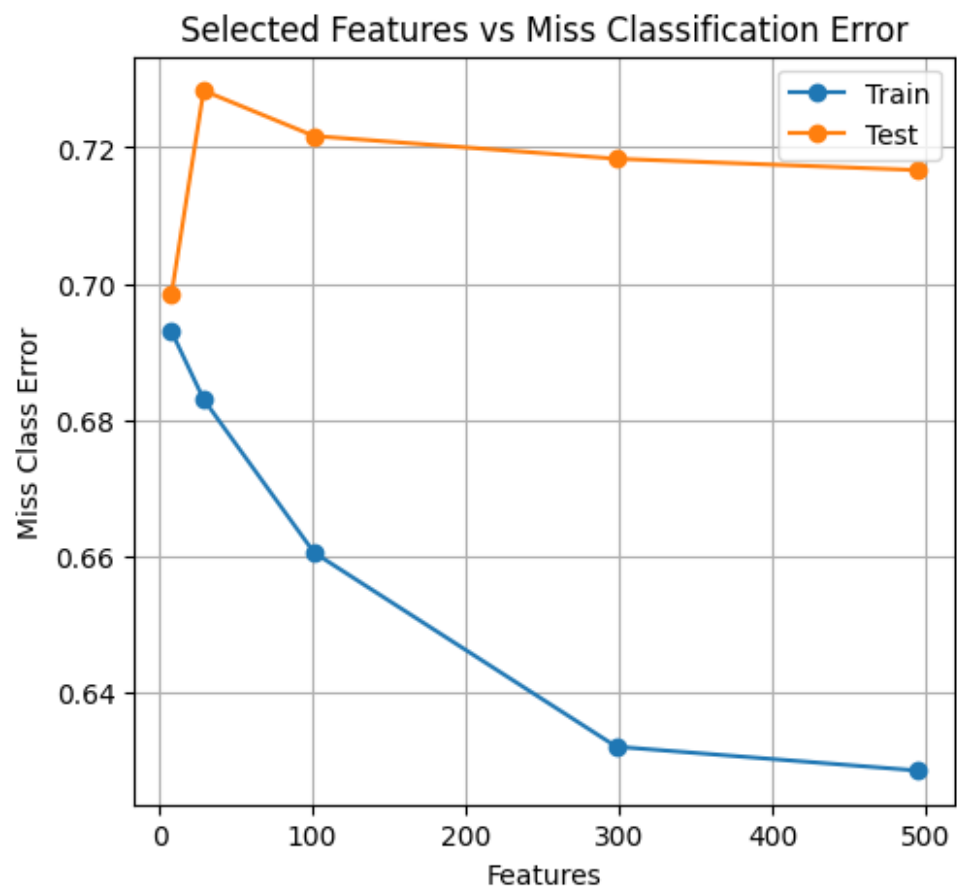
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend()


```



```
plt.show()
```







```
results = pd.DataFrame({  
    'Lambda': thresholds,  
    'Features': features,  
    'Train MisClass Error': train_misclass_errors,  
    'Test MisClass Error': valid_misclass_errors  
})
```

```
print(results)
```

	Lambda	Features	Train MisClass Error	Test MisClass Error
0	0.029990	8	0.6930	0.698333
1	0.024500	29	0.6830	0.728333
2	0.017000	101	0.6605	0.721667
3	0.007500	298	0.6320	0.718333
4	0.000199	494	0.6285	0.716667